

Enhancing Time Series Classification with Self-Supervised Learning

Ali Ismail-Fawaz¹ Maxime Devanne¹ Jonathan Weber¹ and Germain Forestier¹

¹IRIMAS, Université Haute-Alsace, Mulhouse, France

{ali-el-hadi.ismail-fawaz, maxime.devanne, jonathan.weber, germain.forestier}@uha.fr

Keywords: Self Supervised, Time Series Classification, Semi Supervised, Triplet Loss

Abstract: Self-Supervised Learning (SSL) is a range of Machine Learning techniques having the objective to reduce the amount of labeled data required to train a model. In Deep Learning models, SSL is often implemented using specific loss functions relying on pretext tasks leveraging from unlabelled data. In this paper, we explore SSL for the specific task of Time Series Classification (TSC). In the last few years, dozens of Deep Learning architectures were proposed for TSC. However, they almost exclusively rely on the traditional training step involving only labeled data in sufficient numbers. In order to study the potential of SSL for TSC, we propose the **TRIPlet Loss In TimeE (TRILITE)** which relies on an existing triplet loss mechanism and which does not require labeled data. We explore two use cases. In the first one, we evaluate **the interest of TRILITE to boost a supervised classifier when very few labeled data are available**. In the second one, we study the use of TRILITE **in the context of semi-supervised learning, when both labeled and unlabeled data are available**. Experiments performed on 85 datasets from the UCR archive reveal interesting results on some datasets in both use cases.

1 INTRODUCTION

Time Series Classification (TSC) has been a challenging problem addressed by many researchers. While some approaches used basic Machine Learning techniques, more recently, the usage of Deep Learning models has been addressed for classification (Ismail Fawaz et al., 2019; Ismail Fawaz et al., 2020), clustering (Lafabregue et al., 2022; Anowar et al., 2022), averaging (Terefe et al., 2020), adversarial attacks (Pialla et al., 2022b), etc. This is particularly due to the availability of more data, especially after the release of the UCR archive (Dau et al., 2019). The problem at hand is to find a function that maps each sample in the dataset to its corresponding label.

Trying to learn a one-to-one mapping from samples to labels can be sometimes challenging. This is due to the lack of labeled samples a dataset includes. Some approaches use data augmentation, (Fawaz et al., 2018; Pialla et al., 2022a; Kavran et al., 2022) To overcome this problem, Self-Supervised Learning (SSL) offers a way to learn a discriminant latent space without learning to map each sample into a label, which is an advantage over supervised learning.

SSL is used in two main ways, **either using contrastive loss or triplet loss**. On one hand, the contrastive loss **aims to learn how to embed a sample**

and its augmented view to the same point in the latent space. For example, it has been used on image classification (Chen et al., 2020; Grill et al., 2020) and human action recognition (Lin et al., 2020). The triplet loss on the other hand **aims to learn how to embed a sample and its augmented representation (called positive representation) to the same point while learning how to differ those last two from a representation of another sample (called negative representation)**. The triplet loss was introduced in (Schroff et al., 2015) for face recognition and used for knowledge distillation in (Liang et al., 2021; Oki et al., 2020). **The difference between the two losses is the strictness**. The contrastive loss is less strict on choosing what representations should not be close to the original, on contrary to the triplet loss. In other words, the contrastive loss does not consider a negative representation, only positive ones.

Inspired by all those applications, we are interested in adapting the SSL into the time series domain using the triplet loss (Schroff et al., 2015) with our own augmentation method adapted to time series. Although in imaging we have access to large archives such as Imagenet (Deng et al., 2009), it is not the same in the time series domain. In addition, the problem of annotating time series datasets is more frequent than in imaging. For instance, in

imaging, for the majority of datasets, one can use crowd-sourcing to annotate the data samples. This approach is not that easily suitable in the time series domain, given that most of the times, it would be difficult to differentiate between time series samples and annotate them. As a result, for time series, usually one should refer to an expert for the annotation part. To the best of our knowledge, most of the methods that have been developed for SSL on time series evaluate their model considering fully unlabeled data, which is usually not the case in real-world scenarios. In most of these work, SSL is first applied on the data and then a 1-Nearest Neighbor (1-NN) strategy is employed on the output representations for the classification task. However, in real-world scenarios, due to the difficulty to annotate time series, we usually are in one of these two cases: (1) only a small annotated dataset is available or (2) a larger dataset may be available but only a small part of it is annotated. In both of these cases, supervised learning usually overfits quickly because the few amount of annotated samples makes the dataset hard to learn from. To overcome this limitation, SSL is one possible solution by learning representations of time series without the labels information. The question that remains is how to evaluate the self-supervised model?, i.e., how can we know if the model learns meaningful representation of input time series? In this paper we propose a new SSL approach, **TRIPlet Loss In Time** (TRILITE) and evaluate it for these two cases. First, we consider the case of having small annotated time series datasets (**Use case 1**), which we tackle using a self-supervised approach in combination of supervised learning, by concatenating the output representations in order to learn more features. This can be considered as a booster for supervised models if the SSL features are learned in a good way. Second, we consider the case of having larger time series datasets but still with few labeled samples (**Use case 2**), which we tackle using a semi-supervised approach leveraging unlabeled samples. This is to prove that SSL can also take into account unlabeled data in order to learn a more discriminant latent space and produce better representations for the labeled samples.

The main contributions of this papers are:

- a new method of time series augmentation adapted to the context of SSL.
- a new SSL approach based on the triplet loss introduced by (Schroff et al., 2015). Usually other methods use the contrastive loss and we want to assess the triplet loss's contribution in SSL for TSC.
- a new way of evaluating self-supervised model which to the best of our knowledge was not em-

ployed before. Here we showed that with the help of self-supervised combined with supervised learning we can achieve better performance than supervised learning alone on some datasets.

The rest of the paper is organized as follows: in section 2 we introduce some related work on SSL for time series, in section 3 we explain in details our method, in section 4 we evaluate our self-supervised model on the UCR archive (Dau et al., 2019) and we finish with a conclusion in section 5.

2 Related work

2.1 Self-Supervised Learning

SSL for time series was addressed in two main ways, using contrastive learning and triplet loss. It was first addressed for univariate time series in (Franceschi et al., 2019). The authors used a triplet loss to make the model learn how to differentiate between the latent representations of an anchor time series, a positive representation of it and several negative representations. The authors used an unsupervised way to generate their triplets motivated by the problem of variable length time series. While their anchor time series is a sub-sequence of a chosen sample, the positive representation is defined as a sub-sequence of the anchor and the negative representation as a sub-sequence of a randomly selected sample. The advantage of their approach is the insensitivity to variable length time series. Its disadvantage is that it does not perturb the positive representations, hence it facilitates the learning task. This is due to the resemblance between the reference time series and its positive representation.

On the other hand, in (Wickstrøm et al., 2022), the authors addressed the mixing up proposal between two time series. They used a contrastive learning approach to predict the amount of mixing up from two different time series in order to learn representations in the latent space. For each pair of input time series randomly sampled, they create a third sample mixed up from it. This method overcomes the lack of perturbation added to the time series in (Franceschi et al., 2019) but to the best of our knowledge was never tested using a triplet loss.

Originally, in (Schroff et al., 2015), the authors introduced the SSL world with the triplet loss which was motivated from the Siamese network. The authors applied it for face detection using positive and negative examples of the reference image. This induces into the network not only features of the ref-

erence image but also similarities and dissimilarities with other images. This was motivated by the lack of available samples.

Differently, the authors of (Yang et al., 2022) proposed an adaptation of SimCLR (Chen et al., 2020) but with taking into consideration the temporal dependencies between data points in time series.

In addition, some authors contribute in SSL for Multivariate Time Series (MTS), like in (Chen et al., 2022) where the authors constructed a channel-aware self-supervised model using a transformer encoder for MTS. Taking into consideration both channel-wise and time-wise features by training two encoders, one to find the similarity using the contrastive loss and the other to predict the next trend. Another contribution was made in SSL for MTS in (Mohsenvand et al., 2020) where the authors introduced SeqCLR which was an adaptation of SimCLR (Chen et al., 2020). The authors worked on SSL for electroencephalogram classification where they faced two problems. The first was the lack of data and the second was what the type of data augmentation they should use for the contrastive learning.

Distinctively, some work addressed the dimensionality reduction problem like in (Garg, 2021). The authors used an autoencoder where their encoder and decoder were adapted from a LeNet (LeCun et al., 1998) trained on both the reconstruction loss at the output of the decoder and the triplet loss (Schroff et al., 2015) using the triplets generated in the bottleneck layer.

Moreover, the authors in (Eldele et al., 2021) suggested a method based on temporal and contextual contrastive learning using transformers. First they suggested using two different augmentations for the same input time series and feed them to a Siamese network with a FCN encoder. This was followed by a transformer to learn cross-view prediction task of time steps while maximizing similarities between two predictions from different views. This was to help with the forecasting part of the model but at the same time apply contrastive learning on the two predicted representations.

In addition, a review (Lafabregue et al., 2022) on deep representation learning for time series clustering sets the benchmark of clustering techniques on the UCR archive (Dau et al., 2019). The authors varied between which architecture to use, which clustering loss and pretext loss to use. One of the pretext loss in this review was the triplet loss proposed in (Franceschi et al., 2019). The clustering loss is applied on the latent representations.

2.2 Deep learning methods for Time Series Classification

Deep learning for TSC has been a very targeted domain in the past few years. In (Ismail Fawaz et al., 2019), a comparison was made between different architectures starting from Multi Layer Perceptron (MLP) to end up with ResNet (Wang et al., 2017). It was shown in (Ismail Fawaz et al., 2019) how the use of convolution networks like Fully Convolutional Network (FCN) (Wang et al., 2017) and residual connections (ResNet) (He et al., 2016) are significantly better than other approaches. Some work (Mercier et al., 2022) also analyse the explanation of convolution classifiers for time series. More recently, the InceptionTime (Ismail Fawaz et al., 2020) model showed a significant difference in performance by concatenating the output of many convolution layers with different characteristics instead of tuning on one convolution layer. This creates several receptive fields for the model to learn from. In addition, knowledge distillation techniques (Hinton et al., 2015) have been adapted for Time Series Classification (Ay et al., 2022) where the authors transfer the knowledge from a FCN to a smaller convolution neural network.

In our case, we are interested in evaluating the contribution of the triplet loss in the case of TSC. Hence we chose a simple FCN architecture as a backbone of our TRILITE model detailed in the next section.

3 Triplet Loss In Time

3.1 Definitions

Before describing the model, we define some important terms used in the rest of the paper.

Time series. A time series is a sequential data representation of an event changing over time in an equally separated manner. We define a univariate time series as $X = [x_1, x_2, \dots, x_L]$, L being the number of time steps of this time series. A batch $\mathcal{X} = \{\mathbf{X}\}_1^N$ is set of N time series of length L .

Representations. An anchor time series is referred to as *ref*, a positive representation of the anchor as *pos* and a negative representation of the anchor as *neg*.

Metrics. A distance between 2 representations $d(.,.)$ is referred to the Euclidean distance between 2 vectors.




Figure 1: Overview of our TRILITE model.

3.2 Model

Our proposed TRILITE model is composed of **three encoders with weights shared between them**. It corresponds to one encoder fed by the triplets generated as it can be seen in Figure 1. In our case, the encoder architecture is the FCN from (Wang et al., 2017) and implemented in (Ismail Fawaz et al., 2019). However, **we remove the classification layer** as we are considering a self-supervised case. Each triplet’s sample *ref*, *pos* and *neg* are passed through the model to obtain the latent representation ref_l , pos_l and neg_l . These latent representations are of size 128.

3.3 Triplet Loss

Inspired by (Schroff et al., 2015) where triplet loss was used for face recognition, **we propose to use a similar loss for the case of time series**. The triplet loss for a given triplet is defined as:

$$Loss = \max(0, \alpha + d(ref_l, pos_l) - d(ref_l, neg_l)). \quad (1)$$

The goal of this loss is to increase the distance between the ref_l and neg_l but to decrease the distance between ref_l and pos_l . In order **to relax this minimization problem**, an α hyperparameter is added to control the margin between these two distances. As a result we can identify 3 types of triplets as shown in Figure 2:

- **Easy triplet:** The $Loss = 0$ because $d(ref_l, pos_l) + \alpha < d(ref_l, neg_l)$
- **Hard Triplet:** The neg_l is closer to the ref_l than pos_l , $d(ref_l, neg_l) < d(ref_l, pos_l)$
- **Semi-Hard Triplet:** The case is good $d(ref_l, pos_l) < d(ref_l, neg_l)$ but we got a positive loss

We note that if $\alpha = 0$, only two triplets can be defined, the easy triplets and the hard triplets. **Using only easy triplets would overfit the model and using only the hard triplets would underfit the model**. This motivates the use of the hyperparameter α to create




Figure 2: Schema of the relaxed spaced controlled by the margin α

the in-between space of semi-hard triplets. In addition, in 1, the max operation in the loss aims to **transform the problem into a convex one**.

3.4 Triplet generation

In order to generate triplets i.e. *pos* and *neg* generated from *ref*, two main approaches have been proposed. In (Wickstrøm et al., 2022) a mixing up strategy is employed. In (Franceschi et al., 2019) a **masking approach is used**. In this work, our triplet generation approach consists of combining both methods into one, as detailed in Algorithm 1. **First, a *pos* is created using a weighted sum of three time series, the *ref* included**. For the *neg*, the same process is applied except that the *ref* is not included in the weighted sum. We limit the number of mixed up samples to three to ensure that each sample can contribute in a significant manner. This generation process is summarized in the following equations:

$$pos = w * ref + \frac{1-w}{2} * (ts_1 + ts_2) \quad (2)$$

$$neg = w * \bar{ref} + \frac{1-w}{2} * (ts_1 + ts_2) \quad (3)$$

In 2 and 3, ts_1 and ts_2 are randomly selected time series different from the *ref*. Moreover the contribution weight w is randomly selected between 0.6 and 1.0. This ensures the *pos* has more contribution coming from the *ref* than ts_1 and ts_2 .

Second, a mask is generated with a random length and applied on the *pos* and the *neg*. We employed the masking strategy to simplify the training procedure by learning parts of the representations instead of the whole representations.

Third, the unmasked parts of the time series is replaced by a random Gaussian noise. A visualization of the *pos* sample generation can be seen in Figure 3.

We note that during training, the triplet generation is done in an online way at each epoch in order to make the model generalize better.




Figure 3: A *pos* (in orange) is built from three time series including the *ref* (in blue). The resulting time series is close to the *ref* except some areas as highlighted in the red circle. A mask is then applied on the mixed up *pos* to generate the final sample (in green), where the unmasked parts are replaced by a Gaussian noise.

Algorithm 1 Triplet_Generation

```

1: Input: data
2: shuffle(data)
3:  $N \leftarrow \text{data.shape}[0]$   $\triangleright$  Number of samples in data
4:  $l \leftarrow \text{data.shape}[1]$   $\triangleright$  Length of input samples
5:  $w \leftarrow \text{random}(0.6, 1)$   $\triangleright$  The amount of mixing up contribution
6: for  $i : 0 \rightarrow N$  do
7:    $\text{ref}[i] \leftarrow \text{data}[i]$ 
8:    $ts_1 \leftarrow \text{random\_sample}(\text{data})$ 
9:    $ts_2 \leftarrow \text{random\_sample}(\text{data})$ 
10:   $\text{pos}[i] \leftarrow w.\text{ref}[i] + (\frac{1-w}{2}).(ts_1 + ts_2)$ 
11:
12:   $ts_1 \leftarrow \text{random\_sample}(\text{data})$ 
13:   $ts_2 \leftarrow \text{random\_sample}(\text{data})$ 
14:   $ts_2 \leftarrow \text{random\_sample}(\text{data})$ 
15:   $\text{neg}[i] = w.ts_1 + (\frac{1-w}{2}).(ts_2 + ts_3)$ 
16:
17:   $\text{pos}[i], \text{neg}[i] \leftarrow \text{Mask}(\text{pos}[i], \text{neg}[i])$ 
18: end for
19:  $\text{pos} \leftarrow \text{Znormalize}(\text{pos})$ 
20:  $\text{neg} \leftarrow \text{Znormalize}(\text{neg})$ 
21: return  $\text{ref}, \text{pos}, \text{neg}$ 

```

4 Experimental Evaluations

In this section we evaluate the proposed model considering the two use cases identified in Section 1.

4.1 Datasets and implementation details

We based all of our experiments on the UCR archive (Dau et al., 2019) which is the largest open

Algorithm 2 Mask

```

1: Input:  $x, y$ 
2: Output:  $x, y$ 
3:  $l \leftarrow \text{len}(x)$ 
4:  $\text{start} \leftarrow \text{random\_randint}(0, l - 1)$ 
5:  $\text{stop} \leftarrow \text{random\_randint}(\text{start} + \frac{l-1-\text{start}}{10}, \text{start} + \frac{l-1-\text{start}}{2.5})$ 
6:  $x[0 : \text{start}] \leftarrow \text{noise}$ 
7:  $x[\text{stop} + 1 : ] \leftarrow \text{noise}$ 
8:  $y[0 : \text{start}] \leftarrow \text{noise}$ 
9:  $y[\text{stop} + 1 : ] \leftarrow \text{noise}$ 
10: return  $x, y$ 

```

source archive for TSC. We used the 2015 version of the UCR archive made of 85 univariate time series. We z-normalized the datasets before applying our TRILITE model. All of the results are averaged on five runs. We used the Adam optimizer with an initial learning rate of 10^{-3} . For the choice of the hyperparameter α , we tried to tune it on the set of values $\{10^{-6}, 10^{-5}, 10^{-4}, 10^{-3}, 10^{-2}, 10^{-1}\}$. We observed, by visualizing the latent representations, that when we change the value of α the model changes the scale of the representations. By this change of scale, the type of the triplet (as explained in Section 2) would not change. After this observations we fixed the value of α to 10^{-2} tuned on a subset of the UCR archive. We trained the model for 1000 epochs with a batch size of 32. For the evaluation on the test set, we used the last model. We did the experiments on a NVIDIA GeForce GTX 1080 with 8GB of memory. The code is publicly available <https://github.com/MSD-IRIMAS/TRILITE>.

Table 1: Ranking each method on the 85 UCR archive datasets

Method	Rank 1	Rank 2	Rank 3
TRILITE 1-LP	8	4	73
FCN	37	37	11
TRILITE+FCN	40	44	1

4.2 Use case 1: Small annotated time series datasets

To address this first case where a small annotated time series dataset is available, we compared the TRILITE model, followed by a fully connected layer with softmax activation (denoted as TRILITE 1-LP), with the single FCN. The Win-Tie-Loss visualization is reported in Figure 4. We can observe that not surprisingly, the supervised model outperforms the self-supervised one. However, for some datasets we can see that self-supervised features allow to improve the classification accuracy (blue points). This motivated us to evaluate the contribution of self-supervised features in a supervised problem. In order to do that we simply *concatenate* the latent representations of the self-supervised TRILITE model (of size 128) and the latent representations of the supervised FCN model (of size 128) for the train and test sets. Then, these concatenated features are fed to a classifier, a one fully connected layer with a softmax activation. We compare this approach, denoted as TRILITE+FCN, with the single FCN and the single TRILITE models. For each model we compute the classification accuracy and rank them accordingly on each dataset. Ranking results are reported in Table 1 and visualized in the Critical Difference Diagram in Figure 6. These results show that the TRILITE+FCN approach occurs at the first rank position more often than the single FCN model. This is emphasized in the Win-Tie-Loss comparison in Figure 5. We can see that the TRILITE+FCN approach is never worse than the single FCN in a significant manner. This is due to the fact that supervised features can not get perturbed by the SSL features. In the worst case scenario, the linear classifier can learn to reject the SSL features in case they were perturbing the classification. In addition, on average, the difference in accuracy is $3.73 + -9.94$ when TRILITE+FCN wins. Conversely, the difference in accuracy is only $0.64 + -0.7$ when the single FCN wins. This shows that SSL produces features different than the ones produced by supervised learning. As a result, the concatenation of both features allows to improve the classification performance.

This suggests that the self-supervised model is able to discriminate between classes even if it was not its main objective during training. To emphasize this,




Figure 4: TRILITE with 1-LP VS supervised FCN




Figure 5: Concatenation with 1-LP VS supervised FCN

we employed T-distributed Stochastic Neighbor Embedding (TSNE) (Van der Maaten and Hinton, 2008) to visualize a two-dimensional space representing the raw data and the self-supervised features (Figure 7) for the SyntheticControl dataset of the UCR archive. In the later figure, we can clearly identify distinct and compact clusters representing the classes.

4.3 Use case 2: Partially annotated time series datasets

In this second case where, we consider a semi-supervised scenario where only a part of the data is labeled. We aim to evaluate how SSL can overcome this lack of labels. To do that we follow these different steps supposing that only 30% of the training set is labeled:

1. Self-supervised training. We obtain self-supervised latent representations by training our TRILITE model:
 - (a) **experiment 1**: only on the labeled subset.
 - (b) **experiment 2**: on the whole training set.




Figure 6: Average rank mean of the three methods on 84 datasets of the UCR archive.

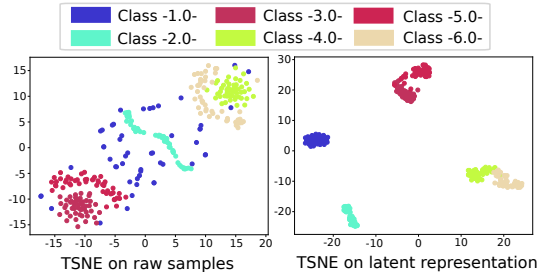


Figure 7: TSNE representation on the SyntheticControl dataset on raw samples and TRILITE latent representation.




Figure 8: Comparison of experiment 1 (1a) and experiment 2 (1b). In experiment 1, the TRILITE model is trained only on the labeled subset (30% of the data). On the contrary, in experiment 2, the TRILITE model is trained on the whole train set. The evaluation is done on the whole test set.

2. Supervised learning. We feed the latent representations of the labeled set (either 1a or 1b) to a Ridge classifier (Peng and Cheng, 2020).
3. Evaluation. The trained classifier is evaluated on the test set.

To not be dependent on a single labeled subset, these steps are repeated over 25 runs and the average accuracy is computed. For each run, the same labeled subset is used for both experiments. The Win-Tie-Loss comparison between experiments 1 and 2 is shown in Figure 8.

We can see that experiment 2 obtains more wins than experiment 1. In addition, on average, the difference in accuracy when experiment 2 wins is $2.12 + -2.13$. On the contrary when experiment 1 wins, the difference in accuracy is on average of $1.17 + -1.21$. This shows that SSL can learn a more meaningful latent representation when taking into consideration the labeled and unlabeled subsets.

5 Conclusion

The difficulty to annotate time series data often results in a lack of labeled data, and thus complicates the training of supervised TSC models. In this paper, we

proposed a Self-Supervised approach for TSC in order to address this problem. In particular, through two use cases, we evaluated how Self-Supervised Learning can be employed in combination with supervised learning to enhance TSC performances. First, we consider the case of small annotated datasets. We showed that with the help of self-supervision, the performance of supervised learning can be improved in some cases. Second, in a case partially annotated datasets, we showed that Self-Supervised Learning can also be a complement to learning supervised models only on labeled data. Hence, these experimental results demonstrated that our SSL approach allows to capture meaningful features that are complementary to features captured in supervised models. The main problem we faced is how to evaluate the self-supervised models, which architecture to use, which data augmentation to apply and which loss to minimize. Furthermore, we aim to study the impact of each hyperparameter on the problem at hand when using Self-Supervised Learning.

ACKNOWLEDGEMENTS

This work was supported by the ANR DELEGATION project (grant ANR-21-CE23-0014) of the French Agence Nationale de la Recherche. The authors would like to acknowledge the High Performance Computing Center of the University of Strasbourg for supporting this work by providing scientific support and access to computing resources. Part of the computing resources were funded by the Equipex Equip@Meso project (Programme Investissements d’Avenir) and the CPER Alsacalcul/Big Data. The authors would also like to thank the creators and providers of the UCR Archive.

REFERENCES

- Anowar, F., Sadaoui, S., and Dalal, H. (2022). Clustering quality of a high-dimensional service monitoring time-series dataset. In *ICAART (2)*, pages 183–192.
- Ay, E., Devanne, M., Weber, J., and Forestier, G. (2022). A study of knowledge distillation in fully convolutional network for time series classification. In *Int. Joint Conference on Neural Networks (IJCNN)*.
- Chen, T., Kornblith, S., Norouzi, M., and Hinton, G. (2020). A simple framework for contrastive learning of visual representations. In *Int. conference on machine learning*, pages 1597–1607. PMLR.
- Chen, Y., Zhou, X., Xing, Z., Liu, Z., and Xu, M. (2022). Cass: A channel-aware self-supervised representation learning framework for multivariate time series classi-

- fication. In *Int. Conference on Database Systems for Advanced Applications*, pages 375–390. Springer.
- Dau, H. A., Bagnall, A., Kamgar, K., Yeh, C.-C. M., Zhu, Y., Gharghabi, S., Ratanamahatana, C. A., and Keogh, E. (2019). The ucr time series archive. *IEEE/CAA Journal of Automatica Sinica*, 6(6):1293–1305.
- Deng, J., Dong, W., Socher, R., Li, L.-J., Li, K., and Fei-Fei, L. (2009). Imagenet: A large-scale hierarchical image database. In *2009 IEEE conference on computer vision and pattern recognition*, pages 248–255. Ieee.
- Eldele, E., Ragab, M., Chen, Z., Wu, M., Kwok, C. K., Li, X., and Guan, C. (2021). Time-series representation learning via temporal and contextual contrasting. *arXiv preprint arXiv:2106.14112*.
- Fawaz, H. I., Forestier, G., Weber, J., Idoumghar, L., and Muller, P.-A. (2018). Data augmentation using synthetic data for time series classification with deep residual networks. *arXiv preprint arXiv:1808.02455*.
- Franceschi, J.-Y., Dieuleveut, A., and Jaggi, M. (2019). Unsupervised scalable representation learning for multivariate time series. *Advances in neural information processing systems*, 32.
- Garg, Y. (2021). Retrim: Reconstructive triplet loss for learning reduced embeddings for multi-variate time series. In *2021 Int. Conference on Data Mining Workshops (ICDMW)*, pages 460–465. IEEE.
- Grill, J.-B., Strub, F., Althé, F., Tallec, C., Richemond, P., Buchatskaya, E., Doersch, C., Avila Pires, B., Guo, Z., Gheshlaghi Azar, M., et al. (2020). Bootstrap your own latent-a new approach to self-supervised learning. *Advances in Neural Information Processing Systems*, 33:21271–21284.
- He, K., Zhang, X., Ren, S., and Sun, J. (2016). Deep residual learning for image recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 770–778.
- Hinton, G., Vinyals, O., Dean, J., et al. (2015). Distilling the knowledge in a neural network.
- Ismail Fawaz, H., Forestier, G., Weber, J., Idoumghar, L., and Muller, P.-A. (2019). Deep learning for time series classification: a review. *Data mining and knowledge discovery*, 33(4):917–963.
- Ismail Fawaz, H., Lucas, B., Forestier, G., Pelletier, C., Schmidt, D. F., Weber, J., Webb, G. I., Idoumghar, L., Muller, P.-A., and Petitjean, F. (2020). Inceptiontime: Finding alexnet for time series classification. *Data Mining and Knowledge Discovery*, 34(6):1936–1962.
- Kavran, D., Zalik, B., and Lukac, N. (2022). Time series augmentation based on beta-vae to improve classification performance. In *ICAART (2)*, pages 15–23.
- Lafabregue, B., Weber, J., Gañarski, P., and Forestier, G. (2022). End-to-end deep representation learning for time series clustering: a comparative study. *Data Mining and Knowledge Discovery*, 36(1):29–81.
- LeCun, Y., Bottou, L., Bengio, Y., and Haffner, P. (1998). Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11):2278–2324.
- Liang, Y., Pan, Y., Lai, H., Liu, W., and Yin, J. (2021). Deep listwise triplet hashing for fine-grained image retrieval. *IEEE Transactions on Image Processing*.
- Lin, L., Song, S., Yang, W., and Liu, J. (2020). Ms2l: Multi-task self-supervised learning for skeleton based action recognition. In *Proceedings of the 28th ACM Int. Conference on Multimedia*, pages 2490–2498.
- Mercier, D., Bhatt, J., Dengel, A., and Ahmed, S. (2022). Time to focus: A comprehensive benchmark using time series attribution methods. *arXiv preprint arXiv:2202.03759*.
- Mohsenvand, M. N., Izadi, M. R., and Maes, P. (2020). Contrastive representation learning for electroencephalogram classification. In *Machine Learning for Health*, pages 238–253. PMLR.
- Oki, H., Abe, M., Miyao, J., and Kurita, T. (2020). Triplet loss for knowledge distillation. In *2020 Int. Joint Conference on Neural Networks (IJCNN)*, pages 1–7. IEEE.
- Peng, C. and Cheng, Q. (2020). Discriminative ridge machine: A classifier for high-dimensional data or imbalanced data. *IEEE Transactions on Neural Networks and Learning Systems*, 32(6):2595–2609.
- Pialla, G., Devanne, M., Weber, J., Idoumghar, L., and Forestier, G. (2022a). Data augmentation for time series classification with deep learning models. In *Advanced Analytics and Learning on Temporal Data (AALTD)*, page undefined. undefined.
- Pialla, G., Fawaz, H. I., Devanne, M., Weber, J., Idoumghar, L., Muller, P.-A., Bergmeir, C., Schmidt, D., Webb, G., and Forestier, G. (2022b). Smooth perturbations for time series adversarial attacks. In *Pacific-Asia Conference on Knowledge Discovery and Data Mining*, pages 485–496. Springer.
- Schroff, F., Kalenichenko, D., and Philbin, J. (2015). Facenet: A unified embedding for face recognition and clustering. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 815–823.
- Terefe, T., Devanne, M., Weber, J., Hailemariam, D., and Forestier, G. (2020). Time series averaging using multi-tasking autoencoder. In *2020 IEEE 32nd Int. Conference on Tools with Artificial Intelligence (IC-TAI)*, pages 1065–1072. IEEE.
- Van der Maaten, L. and Hinton, G. (2008). Visualizing data using t-sne. *Journal of machine learning research*, 9(11).
- Wang, Z., Yan, W., and Oates, T. (2017). Time series classification from scratch with deep neural networks: A strong baseline. In *2017 Int. joint conference on neural networks (IJCNN)*, pages 1578–1585. IEEE.
- Wickstrøm, K., Kampffmeyer, M., Mikalsen, K. Ø., and Jenssen, R. (2022). Mixing up contrastive learning: Self-supervised representation learning for time series. *Pattern Recognition Letters*, 155:54–61.
- Yang, X., Zhang, Z., and Cui, R. (2022). Timeclr: A self-supervised contrastive learning framework for univariate time series representation. *Knowledge-Based Systems*, page 108606.